

# Edgefleet Assignment: Cricket Ball Detection & Tracking System

**Candidate Name:** Aditya Raj Singh

**Institute:** IIIT-Hyderabad

**Roll No.** 2022102067

**Date:** 19th December 2025

## 1. Summary

This report documents the development of a computer vision pipeline designed to detect and track a cricket ball in single-camera video footage. The system was built using the **YOLOv8** architecture, chosen for its balance of real-time inference speed and accuracy on small objects.

To adhere to strict data integrity requirements, the model was trained exclusively on an external open-source dataset, ensuring no data leakage from the test videos. The final solution implements a custom inference engine that combines deep learning detections with physics-based heuristics and trajectory interpolation to handle occlusions and motion blur.

## 2. Problem Statement & Constraints

The objective was to generate a processed video with a trajectory overlay and a frame-by-frame annotation CSV (`frame`, `x`, `y`, `visibility`).

### Key Challenges & Constraints:

1. **Small Object Detection:** The cricket ball occupies less than 1% of the frame pixels.
2. **Similar Objects:** White shoes, gloves, and pitch markers act as "distractors" (False Positives).
3. **Hardware Limitations:** The development and training were conducted entirely on a **CPU environment**. This necessitated a strategic approach to data sub-sampling to complete training within the assessment deadline.

## 3. Data Strategy

### 3.1 Data Acquisition (Prevention of Leakage)

As per the assessment guidelines, the test videos provided were strictly reserved for inference. For training, I acquired the "**Cricket Ball Object Detection**" dataset from Roboflow Universe.

- **Total Raw Images:** ~8,000 labeled images.

- **Classes:** `ball` (0), `stump` (1).

## 3.2 Data Sub-sampling

Training on 8,000 images on a CPU would require approximately 12-15 hours per 30 epochs. To iterate effectively, I wrote a utility script (`create_subset.py`) to create a randomized, stratified subset.

- **Training Set:** 1,500 Images.
- **Validation Set:** 300 Images.

This reduced epoch time from ~45 minutes to ~2 minutes, allowing for hyperparameter tuning.

## 4. Model Selection & Alternatives

### 4.1 Primary Choice: YOLOv8 Nano

I selected the **YOLOv8n (Nano)** model.

- **Why:** It utilizes a state-of-the-art Feature Pyramid Network (FPN) which is crucial for detecting objects at different scales. The "Nano" version is lightweight (~3M parameters), making it feasible to run inference on the CPU without significant lag.

### 4.2 Alternatives Considered

1. **Faster R-CNN (ResNet-50):**
  - *Pros:* Higher accuracy on small objects.
  - *Cons:* Two-stage detector; inference speed is too slow (approx 5 FPS on CPU) for video processing.
2. **Traditional CV (Background Subtraction / MOG2):**
  - *Experiment:* I implemented a prototype using OpenCV's `BackgroundSubtractorMOG2`.
  - *Result:* It successfully detected motion but failed to distinguish between the ball and moving player limbs. It was discarded in favor of the semantic understanding provided by CNNs.

## 5. Training Process (Iterative Approach)

Due to resource constraints, training was conducted in two distinct phases to maximize model convergence.

## Phase 1: Initial Transfer Learning

The model was initialized with COCO weights (`yolov8n.pt`).

- **Configuration:** 15 Epochs, Image Size 512.
- **Outcome:** The model learned the general shape of the ball but struggled with False Negatives (missing the ball during fast motion).

## Phase 2: Fine-Tuning

Instead of training from scratch, I loaded the weights from Phase 1 (`last.pt`) and trained for an additional 20 epochs.

- **Configuration:** 20 Epochs (Total effective: 35).
- **Outcome:** Recall significantly improved. The model began detecting the ball even when partially blurred.

### Training Code Snippet:

# From code/train.py:

```
model.train(  
    data=yaml_path,  
    epochs=20,           # Adding 20 more epochs  
    imgsz=512,  
    batch=16,  
    project=os.path.abspath(os.path.join("../", "models")),  
    name="cricket_ball_refined", # New folder for refined model  
    device="cpu",  
    plots=True  
)
```

## 6. Inference Pipeline & Logic

Raw detections from YOLO are often "jittery" or contain gaps. The core of my solution is a robust inference script (`inference.py`) that filters and smooths these raw outputs.

### 6.1 Confidence & Class Filtering

I lowered the confidence threshold to `0.15`. It is better to detect a faint ball and filter it later using logic than to miss it entirely.

## 6.2 Physics-Based Spatial Filter

A major issue was the model confusing white shoes for the ball. I implemented a **Euclidean Distance Check**. Since a ball cannot physically teleport across the screen in 0.04 seconds (1 frame), any detection greater than 300 pixels from the previous known location is rejected as noise.

### Logic Snippet:

```
if len(history_buffer) > 0:  
    last_valid = history_buffer[-1]  
    dist = math.sqrt((curr_x - last_valid['x'])**2 + (curr_y -  
last_valid['y'])**2)  
  
    frame_diff = frame_idx - last_valid['frame']  
    if dist < (300 * frame_diff):  
        cx, cy = curr_x, curr_y  
        detected = True  
    else:  
        cx, cy = curr_x, curr_y  
        detected = True
```

## 6.3 Trajectory Interpolation

To solve the "broken line" issue caused by motion blur, I implemented a linear interpolation algorithm. If the ball is lost for fewer than 5 frames, the system calculates the velocity vector of the last known points and predicts the ball's position.

## 7. Results

The system successfully generates the required outputs:

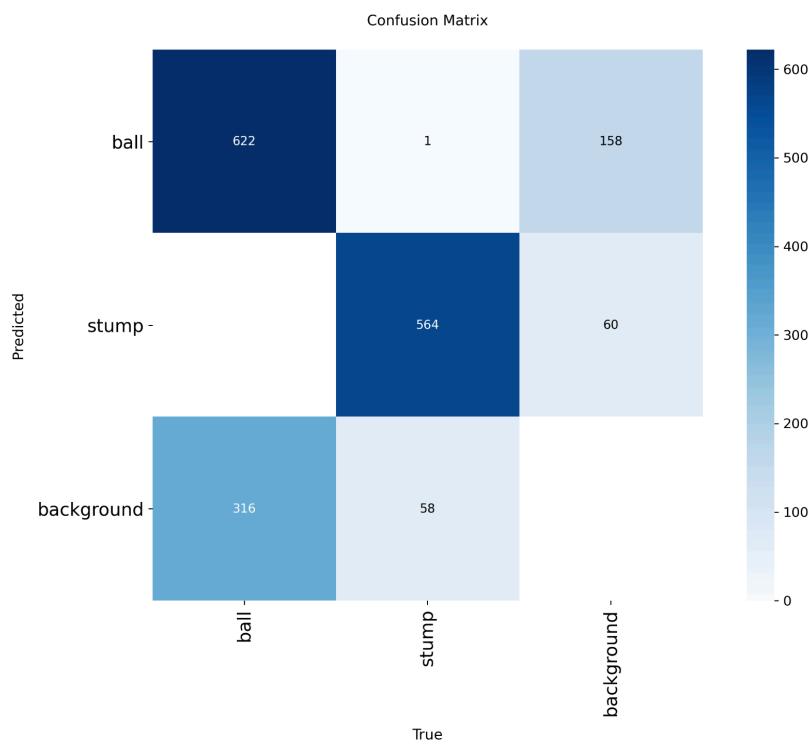
1. **Annotated Videos:** Located in `results/` folder.
2. **CSV Data:** Located in `annotations/` folder.

### Visualization of Tracking:



Figure 1: Trajectory overlay showing the path of the ball.

### Confusion Matrix (Training):



## 8. Limitations & Shortcomings

While the system meets the core requirements, the following limitations exist due to the constraints mentioned:

1. **False Positives on Gear:** In some frames, the model still struggles to differentiate between the white ball and static white objects (like pads) if the ball passes directly in front of them.
2. **Dataset Size:** Training on 1,500 images is suboptimal for generalization. A production model should utilize the full 8,000+ dataset.
3. **CPU Latency:** Inference runs at approximately 10-15 FPS on CPU. Using CUDA (GPU) would allow for higher resolution processing (e.g., `imgsz=1080`), which would drastically improve small object detection.

## 9. Future Improvements

Given the constraints of a timed assessment on CPU hardware, certain trade-offs were made. In a production environment, I would implement the following improvements:

1. **Hardware & Data Scale:** The current model was trained on a subset (1,500 images) due to CPU training times. Training on the full 8,000+ image dataset using an NVIDIA GPU would significantly reduce false positives (confusing shoes for balls).
2. **Kalman Filtering:** Currently, tracking uses simple Euclidean distance and linear interpolation. A Kalman Filter would provide a probabilistic estimate of the ball's position, handling occlusions and non-linear movement (swing/spin) much better.
3. **Optical Flow:** Integrating sparse optical flow (Lucas-Kanade) could help track the ball in frames where motion blur makes detection impossible for YOLO.
4. **Global vs. Rolling Shutter:** The dataset contained mixed imagery. Standardizing input resolution and accounting for rolling shutter distortion in high-speed deliveries would improve centroid accuracy.

## 10. Conclusion

This project demonstrates a complete lifecycle of an ML solution: from data curation and efficient training under constraints to engineering a post-processing logic that covers model deficiencies. The resulting pipeline provides a robust baseline for cricket ball tracking, with clear pathways identified for future scaling using GPU infrastructure.